LEADFLOWAUTOMATION

AN EXECUTIVE BRIEF FOR COMPANY LEADERS

The AI Leverage Blueprint

Where AI belongs in an established company, where it must not go, and how to choose a first move that pays for itself — without betting the company on a demo.

Gergely Racz — PhD in Engineering, University of Cambridge · 10+ years in AI · 7 patents
Architect of production AI systems in daily enterprise use · www.leadflowautomation.net

PART 1

Counsel before tools.

Most AI initiatives in established companies fail the same way: a tool is chosen first, and the company is rearranged around it afterwards. The demo impresses, the pilot starts, and six months later the tool is quietly abandoned — not because the AI was weak, but because nobody had decided what problem it was accountable for solving.

The decision that matters is not **"which AI tool should we buy?"** It is **"which operating constraint, if removed, would create measurable commercial value — and is that constraint actually addressable with today's AI?"** That is a management question, not a technology question, and it has to be answered before any vendor enters the room.

> The order of operations is the whole game: **understand the operation → pick the constraint → then choose (or build) the technology.** Reversing it is how companies end up with impressive demos and unchanged P&Ls.

What "understanding the operation" means in practice

- **Map the real flows** — how a quote, an order, a support case actually moves through the company, including the informal hand-offs no org chart shows.
- **Find where skilled hours pool** — the three processes that quietly consume the most expensive attention every week. Even rough numbers change the conversation.
- **Mark the no-go zones** — decisions that must stay with people, and data that must never reach an external tool. Drawing these lines early is what later makes AI adoption safe.
- **Check the data plumbing** — whether core systems (ERP, CRM, files, mailboxes) can be reached by software at all, or only by manual export. This single fact often decides what is feasible this year versus next.

None of this requires buying anything. All of it determines whether what you eventually buy — or build — will survive contact with your operation.

PART 2

How AI, code, automation and people divide the work.

"AI" is not one thing, and treating it as one thing is expensive. A reliable system almost always combines four kinds of worker, each doing what it is structurally good at:

Deterministic code

Anything with a right answer: calculations, prices, totals, validations, format conversions. Testable, auditable, cheap to run. **If a number matters commercially, code computes it — AI never invents it.**

Automation & integration

Moving data between systems, triggering steps, keeping records in sync. Boring by design. This layer does the repetitive carrying so neither people nor AI have to.

AI models

The genuinely new capability: reading and drafting text, interpreting messy documents, matching and summarising, answering questions over a large body of knowledge. Powerful — and probabilistic, which is why it never gets the final word on consequential actions.

People

Judgment, relationships, accountability. Every consequential action — a quote sent, an order placed, an outbound message — waits for a named person's approval. This is a design principle, not a temporary training wheel.

The architecture that follows from this is simple to state: **AI proposes, code verifies, automation carries, a person approves.** Systems built this way get more useful as trust grows. Systems built the other way — AI acting directly on live operations — get more frightening as usage grows.

> A useful test for any vendor pitch: ask **which of the four roles their product plays**, and what plays the other three. If the answer is "the AI does everything", the missing three layers will become your problem after the invoice clears.

PART 3

Choosing the first operational win.

The first AI project carries a burden no later project will: it sets the company's belief about whether this technology is real. Choose it by three criteria, applied in order:

1. **Commercial value.** The constraint sits on a flow that touches revenue or significant cost — quoting, order handling, supplier communication, customer response times. Not an internal convenience.
2. **Feasibility.** The data the system needs is reachable, the process has a real owner who can say yes, and the ambiguity is bounded. A feasible modest project beats an ambitious stalled one.
3. **Measurability.** You can state, before building, what number should move and by roughly how much — hours saved, response time cut, error rate down. If the payback can't be described in advance, it can't be proven afterwards.

Anti-patterns worth naming

- **The demo trap** — choosing what demos well over what pays. Impressive demos are usually impressive precisely because they skip the integration and approval work that production requires.
- **The moonshot first** — starting with the hardest, most cross-cutting process. The first project should be chosen to be won.
- **The everything platform** — committing to a company-wide "AI transformation" before one process has proven value. Platforms are earned by wins, not announced.

> Phase discipline protects both sides: **each stage is accepted on written evidence before the next is funded.** You decide whether the next stage has been earned — continue, change, pause or stop.

PART 4

Privacy is not a policy document. It's an architecture.

Most companies' real AI exposure today is not the system they might build — it is the tools their people already use. Company information flows into public chatbots with no design, no visibility, and no way to know. The fix is not prohibition (which fails quietly); it is offering a governed alternative built on four structural commitments:

Your tenant, your accounts

Systems run in your cloud environment on your agreements — models reached through your own enterprise endpoints, not a vendor's infrastructure. You can revoke access any afternoon.

Personal data stops at the gate

Where personal data could reach a model, a pseudonymisation layer strips it first — names, addresses, identifiers replaced before any model call, restored only inside your environment. Tested against golden datasets before go-live, logged in production.

Least privilege, human approval

Scoped, read-only access by default. Consequential actions wait for a named person's sign-off — the approval gate is in the architecture, not in a policy PDF.

A paper trail for your lawyer

Data-processing documentation, DPIA support and sub-processor transparency prepared as deliverables — aligned to GDPR and the EU AI Act as it lands.

Built this way, data protection stops being the reason to delay AI and becomes the reason you can adopt it: every boundary is enforced by design, and your data never leaves your control.

PART 5

What this looks like when it's real.

B2B INDUSTRIAL TRADING COMPANY · WESTERN EUROPE

GDPR-gated AI on live commercial data

Problem.
Supplier pricing and quoting consumed skilled hours daily across ERP and mailboxes, and any AI assistance had to be reconciled with strict European data-protection obligations.

Build.
An AI platform in the client's own cloud tenant: supplier-pricing collection and quote preparation integrated with ERP and mail, with a local pseudonymisation gate stripping personal data before any model call — validated against golden test sets to 100% recall at acceptance. Prices are computed in deterministic code; every quote waits for a named person's approval.

Result.
Zero personal records in any model call — by design, validated at acceptance. Each project stage accepted on written evidence before the next was funded.

US ELECTRONICS MANUFACTURER · MULTI-BILLION-DOLLAR REVENUE

Two production RAG systems in daily engineering use

Problem.
Thousands of engineering documents — SOPs, process records, diagrams — held answers the engineering organisation needed daily but could not reach quickly.

Build.
An engineering-SOP retrieval system and a multimodal process-engineering system over 5,000+ documents, text and engineering diagrams. Trust was the design requirement: reranking, mandatory citations on every answer, refusal when confidence is low.

Result.
Both systems in daily production use by the engineering organisation since April 2026.

Engagements are anonymised under contract — the same confidentiality we would apply to yours. Named references available on request where client agreements permit.

PART 6

The leadership questions — and the first 90 days.

Five questions worth asking before any budget is committed. They cost nothing and reorder most AI conversations:

- **Which three processes consume the most skilled hours each week** — and could we put even rough numbers on them?
- **Which decisions must never be delegated to a tool** — and is that written down anywhere?
- **Which company data must never reach an external AI** — and would anyone currently know if it did?
- **Can our core systems be reached by software**, or only by manual export?
- **What evidence would justify the second project?** Define it before starting the first.

Days 1–30	Map the operation: flows, skilled-hour pools, data reachability, no-go zones. Write the one-page interim guideline on public AI tools. Output: a short written map of where AI belongs — and where it must not go.
Days 31–60	Select the first win by value × feasibility × measurability. Define its acceptance evidence and baseline numbers. Scope the pseudonymisation and approval gates it needs.
Days 61–90	Build the first controlled slice — working software demonstrated on real data behind human approval. Measure against the baseline. Decide, on evidence, whether stage two is earned.

Ninety days is enough to know — not to finish, but to know whether AI has a commercial place in your operation and what it is. That knowledge, written down, is worth more than most pilots.

THE NEXT STEP, IF YOU WANT IT

Your free AI Operating Map.

This brief gives you the framework. The Operating Map applies it to your company specifically: where AI belongs in **your** operation, where it must not go, and the one first move worth funding.

- **One 30-minute working session** — with the engineer who would do the work, not a sales team.
- **A written map within 5 business days** — where AI creates leverage in your company, where it must not go, which systems and data are involved, and the one controlled first win worth funding.
- **No obligation.** If nothing there is worth building, that is a useful answer too — and you keep the map.

Book the 30-minute session

Scan the QR on the banner, or go to:

www.leadflowautomation.net/services

About the author

Gergely Racz holds a PhD in Engineering from the University of Cambridge, has worked in AI and machine learning for over ten years, and holds seven patents. He designs and builds AI systems for European companies — from GDPR-gated agent platforms on live commercial data to production retrieval systems in daily use at a multi-billion-dollar US manufacturer — and builds every engagement personally.